

Assignment 2: Lenia — An Artificial Life

Plabon Shaha and Guillem Masdemont

University of Ljubljana, Spring Semester, EMAI Cohort 25–27

April 14, 2026

Exercise 1: We parallelize the algorithm and present 5 versions that keep optimizing. We compare the speedup performance with the baseline model. We consider the following 5 models:

- **GPU Naive:** An embarrassingly parallel implementation on CUDA by blocks (see `heat1.cu`).
- **GPU Shared Memory (V1):** We use `__shared__` memory to allocate the halo pixels that are not inside the block in the computation (see `heat2.cu` and `heat3.cu`).
- **GPU V2 (Constant + No Modulo):** Relocates the kernel to `__constant__` memory for hardware broadcasting and replaces expensive modulo (%) operations with `if/else` bounds checking to reduce arithmetic costs.
- **GPU V3 (FP32):** Transitions from 64-bit to 32-bit precision. By this, we unlock the GPU's dominant FP32 cores and doubles effective memory bandwidth.
- **GPU V4 (Fused + Texture):** Combines multiple memory accesses into a single operation and utilizes texture memory for spatial locality, further enhancing performance.




Figure 1: Performance comparison of CUDA optimization layers. **Left:** Log-log plot showing execution time versus grid size, highlighting the algebraic scaling advantage of the GPU implementations. **Right:** Grouped bar chart detailing the speedup factor of each GPU kernel relative to the sequential CPU baseline for large grid sizes.

Now, we plot the relative speed performance as a function of the image size and the tile size: The staggered implementation of CUDA optimizations resulted in a peak speedup of nearly 7700 $\times$ over the sequential CPU baseline. The impact of each architectural consideration can be summarized as follows:




Figure 2: Comparison of the overall speedup factor across different thread block tile sizes.

Exercise 2: We now present results for part 2. In this section, we compare an improved CPU baseline against advanced GPU variants. We consider the following implementations:

- **CPU Sequential + OpenMP** (`evolve_lenia_seq_v2`): Refactors the original CPU baseline and parallelizes dominant loops with OpenMP. This version is used as the reference T_{seq} for speedup calculations.
- **GPU V5 (Autotuned FP32, `evolve_lenia_v5`)**: Evaluates several CUDA thread-block geometries at startup and automatically selects the fastest configuration for the target GPU, then runs the full FP32 simulation with that tile shape.
- **GPU V6 (Dual-GPU Domain Split, `evolve_lenia_v6`)**: Splits the world horizontally across two GPUs and exchanges halo rows every step (peer-to-peer when available, otherwise host-staged), enabling larger throughput than single-device execution.

Appendix: Recollected data

Table 1: Performance Metrics with Uncertainty (Tile Size 8)

N	T_{seq}	T_{naive}	T_{v1}	T_{v2}	$T_{v3}(f)$	$T_{v4}(f)$	Speedup(v4)
128	5.0847	0.0465 ± 0.0983	0.0049 ± 0.0007	0.0053 ± 0.0040	0.0031 ± 0.0001	0.0035 ± 0.0009	1472.12
256	20.3207	0.0171 ± 0.0008	0.0098 ± 0.0005	0.0077 ± 0.0003	0.0065 ± 0.0003	0.0064 ± 0.0003	3153.43
512	81.3130	0.0501 ± 0.0035	0.0330 ± 0.0015	0.0283 ± 0.0095	0.0223 ± 0.0003	0.0218 ± 0.0005	3723.12
1024	325.2871	0.1750 ± 0.0014	0.1309 ± 0.0155	0.0873 ± 0.0018	0.0828 ± 0.0004	0.0803 ± 0.0009	4052.11
2048	1301.3061	0.6970 ± 0.0009	0.4971 ± 0.0011	0.3512 ± 0.0008	0.3206 ± 0.0003	0.3125 ± 0.0002	4163.95
4096	5205.5831	2.7612 ± 0.0020	1.9561 ± 0.0059	1.3998 ± 0.0025	1.2719 ± 0.0001	1.2382 ± 0.0001	4204.24

Table 2: Performance Metrics with Uncertainty (Tile Size 16)

N	T_{seq}	T_{naive}	T_{v1}	T_{v2}	$T_{v3}(f)$	$T_{v4}(f)$	Speedup(v4)
128	5.0795	0.0466 ± 0.0961	0.0094 ± 0.0152	0.0037 ± 0.0001	0.0028 ± 0.0001	0.0030 ± 0.0004	1670.89
256	20.3176	0.0170 ± 0.0008	0.0076 ± 0.0003	0.0053 ± 0.0007	0.0050 ± 0.0019	0.0045 ± 0.0001	4557.75
512	81.2717	0.0500 ± 0.0057	0.0239 ± 0.0027	0.0155 ± 0.0007	0.0128 ± 0.0013	0.0123 ± 0.0011	6610.68
1024	325.1207	0.1760 ± 0.0072	0.0890 ± 0.0061	0.0744 ± 0.0226	0.0454 ± 0.0012	0.0445 ± 0.0029	7322.62
2048	1300.6583	0.6915 ± 0.0103	0.3539 ± 0.0108	0.2934 ± 0.0207	0.1729 ± 0.0068	0.1732 ± 0.0200	7513.91
4096	5203.7136	2.7695 ± 0.0163	1.3920 ± 0.0069	1.2082 ± 0.0033	0.6839 ± 0.0028	0.6615 ± 0.0035	7866.53

Table 3: Performance Metrics with Uncertainty (Tile Size 32)

N	T_{seq}	T_{naive}	T_{v1}	T_{v2}	$T_{v3}(f)$	$T_{v4}(f)$	Speedup(v4)
128	5.1617	0.0278 ± 0.0470	0.0081 ± 0.0005	0.0053 ± 0.0003	0.0049 ± 0.0014	0.0049 ± 0.0012	1058.81
256	20.3258	0.0191 ± 0.0060	0.0077 ± 0.0015	0.0048 ± 0.0001	0.0040 ± 0.0001	0.0040 ± 0.0001	5081.45
512	81.2846	0.0478 ± 0.0019	0.0271 ± 0.0011	0.0189 ± 0.0078	0.0150 ± 0.0020	0.0153 ± 0.0046	5314.11
1024	325.0540	0.1730 ± 0.0014	0.0928 ± 0.0114	0.0585 ± 0.0114	0.0489 ± 0.0019	0.0478 ± 0.0057	6793.61
2048	1300.5371	0.6897 ± 0.0022	0.3611 ± 0.0023	0.2248 ± 0.0011	0.1843 ± 0.0014	0.1742 ± 0.0003	7464.55
4096	5205.6769	2.7613 ± 0.0105	1.4221 ± 0.0083	0.9046 ± 0.0065	0.7161 ± 0.0080	0.6881 ± 0.0083	7564.73